

Autoencoders Lab

visualizing HapMap phase 3 populations

This is a solution template. Every chunk of code requiring your input will begin with the # TASK comment and all places where you should fill-in with your code are marked by ellipsis (...).

Stage 0 — getting the data

First, we will download the data from the linked Dropbox account. The code is hidden as it is not super important here. Double-click below if you are curious to see it!

```
In [ ]: # Small config to select the GPU in our local machine
import tensorflow as tf

# GPU list
gpus = tf.config.list_physical_devices('GPU')

# Select GPU 1
tf.config.set_visible_devices(gpus[1], 'GPU')
```

```
In [ ]: #@title Load raw data. Double click to see the code. { display-mode: "form" }

!wget https://www.dropbox.com/s/g7862q1l4ls9z3x/autosomal_5k_matrix.csv
!wget https://www.dropbox.com/s/3lv0062dw20qdqg/autosomal_5k_phenos.csv
!wget https://www.dropbox.com/s/6nzsrxkm536a5j/autosomal_5k_kinship.csv
```

Now, we will load the data and make sure they look as expected. Note, the genotypes per individual (row) are encoded as the count of minor alleles and thus can take values `gt = {0, 1, 2}`.

```
In [2]: import numpy as np
import pandas as pd
import tensorflow as tf
import seaborn as sns
from tensorflow import keras
from tensorflow.keras import layers
from matplotlib import pyplot as plt

data = pd.read_csv("autosomal_5k_matrix.csv", header=0, index_col=0)
pheno = pd.read_csv("autosomal_5k_phenos.csv", header=0, index_col=0)

print(data)
print(pheno)
```

	rs416967	rs17013842	rs13052452	rs11049986	rs10994341	rs1504289	\	
NA19919	2	1	0	1	1	0		
NA19916	1	0	0	1	1	0		
NA19835	0	0	1	1	0	1		
NA20282	1	1	0	1	0	0		
NA19703	0	0	0	0	1	1		
...		
NA19119	1	0	1	0	1	1		
NA18860	1	0	0	1	1	0		
NA19207	2	0	0	0	1	0		
NA19103	1	0	1	0	0	0		
NA19099	0	0	1	1	0	1		
	rs882529	rs3885937	rs537330	rs9372090	...	rs1558766	rs7818288	\
NA19919	0	1	1	1	...	0	1	
NA19916	1	1	2	1	...	0	0	
NA19835	0	2	2	1	...	0	1	
NA20282	0	0	1	1	...	0	0	
NA19703	0	0	0	2	...	0	1	
...	
NA19119	1	1	1	2	...	0	0	
NA18860	0	1	2	1	...	1	1	
NA19207	1	1	2	0	...	0	0	
NA19103	1	1	2	1	...	0	0	
NA19099	0	1	0	0	...	0	1	
	rs1051685	rs11223492	rs789492	rs6557516	rs7313246	rs317892	\	
NA19919	0	0	1	0	1	1		
NA19916	0	0	0	1	1	1		
NA19835	0	0	0	0	2	2		
NA20282	0	0	0	0	1	2		
NA19703	1	0	1	1	1	1		
...		
NA19119	0	0	0	0	0	1		
NA18860	0	0	0	0	1	0		
NA19207	0	0	0	1	0	1		
NA19103	0	0	0	0	2	1		
NA19099	0	0	0	0	1	1		
	rs11937009	rs2806497						
NA19919	1	0						
NA19916	1	0						
NA19835	0	0						
NA20282	0	0						
NA19703	2	0						
...						
NA19119	0	0						
NA18860	0	1						
NA19207	1	0						
NA19103	2	0						
NA19099	0	0						

[1184 rows x 5000 columns]

	id	sex	FID	dad	mom	pheno	population
NA19919	NA19919	1	2427	NA19908	NA19909	0	ASW
NA19916	NA19916	1	2431	0	0	0	ASW
NA19835	NA19835	0	2424	0	0	0	ASW
NA20282	NA20282	0	2469	0	0	0	ASW
NA19703	NA19703	1	2368	0	0	0	ASW
...	
NA19119	NA19119	1	Y060	0	0	0	YRI
NA18860	NA18860	1	Y012	NA18859	NA18858	0	YRI
NA19207	NA19207	1	Y051	0	0	0	YRI
NA19103	NA19103	1	Y042	NA19101	NA19102	0	YRI
NA19099	NA19099	0	Y105	0	0	0	YRI

[1184 rows x 7 columns]

Now, we will create a dictionary and re-name our populations so that the names are a bit more informative:

```
In [3]: pop_dict = {'ASW':'African ancestry in SW USA',
                 'CEU':'Utah residents with N and W European ancestry',
                 'CHB':'Han Chinese in Beijing China',
                 'CHD':'Chinese in Metropolitan Denver Colorado',
                 'GIH':'Gujarati Indians in Houston Texas',
                 'JPT':'Japanese in Tokyo Japan',
                 'LWK':'Luhya in Webuye Kenya',
                 'MEX':'Mexican ancestry in Los Angeles California',
                 'MKK':'Maasai in Kinyawa Kenya',
                 'TSI':'Toscans in Italy',
                 'YRI':'Yoruba in Ibadan Nigeria'}
pheno2 = pheno.replace({"population": pop_dict})
pheno2
```

Out[3]:

	id	sex	FID	dad	mom	pheno	population
NA19919	NA19919	1	2427	NA19908	NA19909	0	African ancestry in SW USA
NA19916	NA19916	1	2431	0	0	0	African ancestry in SW USA
NA19835	NA19835	0	2424	0	0	0	African ancestry in SW USA
NA20282	NA20282	0	2469	0	0	0	African ancestry in SW USA
NA19703	NA19703	1	2368	0	0	0	African ancestry in SW USA
...
NA19119	NA19119	1	Y060	0	0	0	Yoruba in Ibadan Nigeria
NA18860	NA18860	1	Y012	NA18859	NA18858	0	Yoruba in Ibadan Nigeria
NA19207	NA19207	1	Y051	0	0	0	Yoruba in Ibadan Nigeria
NA19103	NA19103	1	Y042	NA19101	NA19102	0	Yoruba in Ibadan Nigeria
NA19099	NA19099	0	Y105	0	0	0	Yoruba in Ibadan Nigeria

1184 rows × 7 columns

```
In [4]: # === TASK 1 ===
print("Example loci / marker names:")
print(data.columns[:10])

print("Total number of individuals:", len(pheno2))

print("\nNumber of individuals per population:")
print(pheno2['population'].value_counts())

print("\nSex distribution (1 = male, 0 = female):")
print(pheno2['sex'].value_counts())

print("\nSex distribution per population:")
print(pheno2.groupby('population')['sex'].value_counts())

# Family structure
family_sizes = pheno2['FID'].value_counts()
num_families = family_sizes.shape[0]

print("\nTotal number of families:", num_families)
print("\nFamily size statistics:")
print(family_sizes.describe())

# Estimation of siblings structure
full_sibs = family_sizes[family_sizes >= 3].count()
half_sibs = family_sizes[family_sizes == 2].count()

print("\nEstimated number of families with full-siblings (size >= 3):", full_sibs)
print("Estimated number of families with half-siblings (size = 2):", half_sibs)
```

```
Example loci / marker names:  
Index(['rs416967', 'rs17013842', 'rs13052452', 'rs11049986', 'rs10994341',  
      'rs1504289', 'rs882529', 'rs3885937', 'rs537330', 'rs9372090'],  
      dtype='object')  
Total number of individuals: 1184
```

Number of individuals per population:

population	
Maasai in Kinyawa Kenya	171
Yoruba in Ibadan Nigeria	167
Utah residents with N and W European ancestry	165
Luhya in Webuye Kenya	90
Gujarati Indians in Houston Texas	88
Toscans in Italy	88
Japanese in Tokyo Japan	86
Chinese in Metropolitan Denver Colorado	85
Han Chinese in Beijing China	84
African ancestry in SW USA	83
Mexican ancestry in Los Angeles California	77

Name: count, dtype: int64

Sex distribution (1 = male, 0 = female):

sex	
0	595
1	589

Name: count, dtype: int64

Sex distribution per population:

population	sex	
African ancestry in SW USA	0	45
	1	38
Chinese in Metropolitan Denver Colorado	0	44
	1	41
Gujarati Indians in Houston Texas	1	45
	0	43
Han Chinese in Beijing China	0	42
	1	42
Japanese in Tokyo Japan	1	44
	0	42
Luhya in Webuye Kenya	0	45
	1	45
Maasai in Kinyawa Kenya	1	86
	0	85
Mexican ancestry in Los Angeles California	0	43
	1	34
Toscans in Italy	0	44
	1	44
Utah residents with N and W European ancestry	0	85
	1	80
Yoruba in Ibadan Nigeria	1	90
	0	77

Name: count, dtype: int64

Total number of families: 770

Family size statistics:

count	770.000000
mean	1.537662
std	1.033514
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	6.000000

Name: count, dtype: float64

Estimated number of families with full-siblings (size >= 3): 147

Estimated number of families with half-siblings (size = 2): 60

```
In [5]: # TASK Scaling  
# We need to scale our counts data so that it is bound between 0 and 1.
```

```

from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
geno_data = pd.DataFrame(scaler.fit_transform(data), index=data.index, columns=data.columns)
geno_data

```

Out[5]:

	rs416967	rs17013842	rs13052452	rs11049986	rs10994341	rs1504289	rs882529	rs388591
NA19919	1.0	0.5	0.0	0.5	0.5	0.0	0.0	0.0
NA19916	0.5	0.0	0.0	0.5	0.5	0.0	0.5	0.0
NA19835	0.0	0.0	0.5	0.5	0.0	0.5	0.0	0.0
NA20282	0.5	0.5	0.0	0.5	0.0	0.0	0.0	0.0
NA19703	0.0	0.0	0.0	0.0	0.5	0.5	0.0	0.0
...
NA19119	0.5	0.0	0.5	0.0	0.5	0.5	0.5	0.0
NA18860	0.5	0.0	0.0	0.5	0.5	0.0	0.0	0.0
NA19207	1.0	0.0	0.0	0.0	0.5	0.0	0.5	0.0
NA19103	0.5	0.0	0.5	0.0	0.0	0.0	0.5	0.0
NA19099	0.0	0.0	0.5	0.5	0.0	0.5	0.0	0.0

1184 rows × 5000 columns

In [6]:

```

# TASK Randomly split into the training and the validation set, so that 80 per-cent of individuals are in the training set

train = geno_data.sample(frac=0.8, random_state=42) # frac=0.8 → 80% training
test = geno_data.drop(train.index)
train = train.reset_index(drop=True)
test = test.reset_index(drop=True)

# TASK Print some info about the resulting split
print("Total number of individuals:", len(geno_data))
print("\t - training set:", len(train))
print("\t - test set:", len(test))

```

Total number of individuals: 1184

- training set: 947
- test set: 237

In [7]:

```

# TASK Specify the autoencoder model

train_tensor = train.to_numpy()
print(type(train_tensor))

# TASK Hyperparameters
# use ReLu activations, ADAM optimizer and
# mean squared error as the loss function
hp_loss_fn = 'mse'
hp_act_fn = 'relu'
hp_optimizer = 'adam'
hp_metrics = ['mse', 'mae', 'mape']

input_data = keras.Input(shape = (train_tensor.shape[1],))

# TASK Define architecture of the encoder:
# the second layer should be a batch normalization

# Baseline
def Encoder_bl(input):
    # Encoder
    layer1 = layers.Dense(units = 512, activation = hp_act_fn, name='layer1')(input)
    layer2 = layers.BatchNormalization(name='layer2')(layer1) # batch norm
    layer3 = layers.Dropout(rate = 0.05, name='layer3')(layer2)
    layer4 = layers.Dense(units = 128, activation = hp_act_fn, name='layer4')(layer3)

```

```

layer5 = layers.Dropout(rate = 0.025, name='layer5')(layer4)
layer6 = layers.Dense(units = 25, activation = hp_act_fn, name='layer6')(layer5)
bottleneck = layers.Dense(units = 2, name='layer_bottleneck')(layer6)
return(bottleneck)

# TASK Look at the encoder, complete the decoder function

def Decoder_bl(bottleneck):
    # Decoder
    layer7 = layers.Dense(units = 128, activation = hp_act_fn)(bottleneck)
    layer8 = layers.Dropout(rate=0.025, name='layer8')(layer7)
    layer9 = layers.Dense(units=512, activation=hp_act_fn, name='layer9')(layer8)
    layer10 = layers.Dropout(rate=0.05, name='layer10')(layer9)
    layer11 = layers.Dense(units=1500, activation=hp_act_fn, name='layer11')(layer10)
    layer12 = layers.Dense(units = train_tensor.shape[1], activation = 'sigmoid')(layer11)
    return(layer12)

def Autoencoder_bl(input):
    enc = Encoder_bl(input)
    autoenc = Decoder_bl(enc)
    return(autoenc)

# Smaller

def Encoder_sm(input):
    # Encoder
    layer1 = layers.Dense(units = 512, activation = hp_act_fn, name='layer1')(input)
    layer2 = layers.BatchNormalization(name='layer2')(layer1) # batch norm
    layer3 = layers.Dropout(rate = 0.05, name='layer3')(layer2)
    layer4 = layers.Dense(units = 128, activation = hp_act_fn, name='layer4')(layer3)
    bottleneck = layers.Dense(units = 2, name='layer_bottleneck')(layer4)
    return(bottleneck)

def Decoder_sm(bottleneck):
    # Decoder
    layer7 = layers.Dense(units = 128, activation = hp_act_fn)(bottleneck)
    layer8 = layers.Dropout(rate=0.025, name='layer8')(layer7)
    layer9 = layers.Dense(units=512, activation=hp_act_fn, name='layer9')(layer8)
    layer10 = layers.Dense(units = train_tensor.shape[1], activation = 'sigmoid')(layer9)
    return(layer10)

def Autoencoder_sm(input):
    enc = Encoder_sm(input)
    autoenc = Decoder_sm(enc)
    return(autoenc)

#Bigger

def Encoder_bg(input):
    # Encoder
    layer1 = layers.Dense(units = 2000, activation = hp_act_fn, name='layer1')(input)
    layer2 = layers.BatchNormalization(name='layer2')(layer1) # batch norm
    layer3 = layers.Dropout(rate = 0.05, name='layer3')(layer2)
    layer4 = layers.Dense(units = 500, activation = hp_act_fn, name='layer4')(layer3)
    layer5 = layers.Dropout(rate = 0.025, name='layer5')(layer4)
    layer6 = layers.Dense(units = 50, activation = hp_act_fn, name='layer6')(layer5)
    bottleneck = layers.Dense(units = 2, name='layer_bottleneck')(layer6)
    return(bottleneck)

def Decoder_bg(bottleneck):
    # Decoder
    layer7 = layers.Dense(units = 50, activation = hp_act_fn)(bottleneck)
    layer8 = layers.Dropout(rate=0.025, name='layer8')(layer7)
    layer9 = layers.Dense(units=500, activation=hp_act_fn, name='layer9')(layer8)
    layer10 = layers.Dropout(rate=0.05, name='layer10')(layer9)
    layer11 = layers.Dense(units=2000, activation=hp_act_fn, name='layer11')(layer10)
    layer12 = layers.Dense(units = train_tensor.shape[1], activation = 'sigmoid')(layer11)
    return(layer12)

def Autoencoder_bg(input):
    enc = Encoder_bg(input)
    autoenc = Decoder_bg(enc)
    return(autoenc)

```

```

autoencoder_model_bl = keras.Model(inputs = input_data, outputs = Autoencoder_bl(input_data))
autoencoder_model_bl.compile(
    loss = hp_loss_fn,
    optimizer = hp_optimizer,
    metrics = hp_metrics
)

autoencoder_model_sm = keras.Model(inputs = input_data, outputs = Autoencoder_sm(input_data))
autoencoder_model_sm.compile(
    loss = hp_loss_fn,
    optimizer = hp_optimizer,
    metrics = hp_metrics
)
autoencoder_model_bg = keras.Model(inputs = input_data, outputs = Autoencoder_bg(input_data))
autoencoder_model_bg.compile(
    loss = hp_loss_fn,
    optimizer = hp_optimizer,
    metrics = hp_metrics
)

# TASK Visualise the created architecture and summarise its parameters
autoencoder_model_bl.summary()
autoencoder_model_sm.summary()
autoencoder_model_bg.summary()

```

<class 'numpy.ndarray'>

I0000 00:00:1764239057.822997 2332417 gpu_device.cc:2019] Created device /job:localhost/replica
a:0/task:0/device:GPU:0 with 46761 MB memory: -> device: 1, name: NVIDIA RTX A6000, pci bus i
d: 0000:b3:00.0, compute capability: 8.6

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 5000)	0
layer1 (Dense)	(None, 512)	2,560,512
layer2 (BatchNormalization)	(None, 512)	2,048
layer3 (Dropout)	(None, 512)	0
layer4 (Dense)	(None, 128)	65,664
layer5 (Dropout)	(None, 128)	0
layer6 (Dense)	(None, 25)	3,225
layer_bottleneck (Dense)	(None, 2)	52
dense (Dense)	(None, 128)	384
layer8 (Dropout)	(None, 128)	0
layer9 (Dense)	(None, 512)	66,048
layer10 (Dropout)	(None, 512)	0
layer11 (Dense)	(None, 1500)	769,500
dense_1 (Dense)	(None, 5000)	7,505,000

Total params: 10,972,433 (41.86 MB)

Trainable params: 10,971,409 (41.85 MB)

Non-trainable params: 1,024 (4.00 KB)

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 5000)	0
layer1 (Dense)	(None, 512)	2,560,512
layer2 (BatchNormalization)	(None, 512)	2,048
layer3 (Dropout)	(None, 512)	0
layer4 (Dense)	(None, 128)	65,664
layer_bottleneck (Dense)	(None, 2)	258
dense_2 (Dense)	(None, 128)	384
layer8 (Dropout)	(None, 128)	0
layer9 (Dense)	(None, 512)	66,048
dense_3 (Dense)	(None, 5000)	2,565,000

Total params: 5,259,914 (20.06 MB)

Trainable params: 5,258,890 (20.06 MB)

Non-trainable params: 1,024 (4.00 KB)

Model: "functional_2"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 5000)	0
layer1 (Dense)	(None, 2000)	10,002,000
layer2 (BatchNormalization)	(None, 2000)	8,000
layer3 (Dropout)	(None, 2000)	0
layer4 (Dense)	(None, 500)	1,000,500
layer5 (Dropout)	(None, 500)	0
layer6 (Dense)	(None, 50)	25,050
layer_bottleneck (Dense)	(None, 2)	102
dense_4 (Dense)	(None, 50)	150
layer8 (Dropout)	(None, 50)	0
layer9 (Dense)	(None, 500)	25,500
layer10 (Dropout)	(None, 500)	0
layer11 (Dense)	(None, 2000)	1,002,000
dense_5 (Dense)	(None, 5000)	10,005,000

Total params: 22,068,302 (84.18 MB)

Trainable params: 22,064,302 (84.17 MB)

Non-trainable params: 4,000 (15.62 KB)

```
In [ ]: # TASK Set hyperparameters for model fitting
# Begin by training for 30 epochs, with mini-batch of 256 and validation set
# having 20 per-cent of examples

hp_epochs = 30
hp_batch_size = 256      # mini-batch size
hp_val_split = 0.2       # 20% of training data for validation

Autoencoder_bl = autoencoder_model_bl.fit(x = train_tensor,
```

```

y = train_tensor,
epochs = hp_epochs,
batch_size = hp_batch_size,
shuffle = True,
validation_split = hp_val_split,
verbose=0
)

Autoencoder_sm = autoencoder_model_sm.fit(x = train_tensor,
                                           y = train_tensor,
                                           epochs = hp_epochs,
                                           batch_size = hp_batch_size,
                                           shuffle = True,
                                           validation_split = hp_val_split,
                                           verbose=0
                                         )

Autoencoder_bg = autoencoder_model_bg.fit(x = train_tensor,
                                           y = train_tensor,
                                           epochs = hp_epochs,
                                           batch_size = hp_batch_size,
                                           shuffle = True,
                                           validation_split = hp_val_split,
                                           verbose=0
                                         )

```

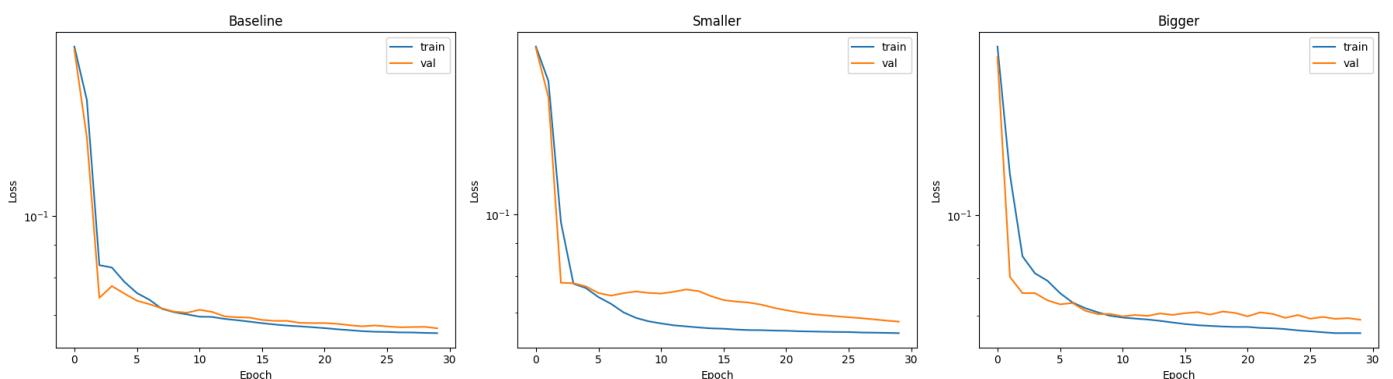
Now, let us look at the training process:

```
In [9]: histories = [Autoencoder_bl, Autoencoder_sm, Autoencoder_bg]
titles = ['Baseline', 'Smaller', 'Bigger']

fig, axs = plt.subplots(1, 3, figsize=(18, 5))

for i, history in enumerate(histories):
    axs[i].plot(history.history['loss'], label='train')
    axs[i].plot(history.history['val_loss'], label='val')
    axs[i].set_title(titles[i])
    axs[i].set_xlabel('Epoch')
    axs[i].set_ylabel('Loss')
    axs[i].set_yscale('log')
    axs[i].legend(loc='upper right')

plt.tight_layout()
plt.show()
```



Now, that the model is trained, we can save the weights and use them to build an encoder. Note that weights are saved for the entire autoencoder, so we need to use `skip_mismatch = True` along with `by_name = True` to initialize weights in our encoder.

```
In [10]: autoencoder_model_bl.save_weights('autoencoder_bl_weights.weights.h5',
                                         overwrite = True)

encoder_model_bl = keras.Model(inputs = input_data, outputs = Encoder_bl(input_data))
encoder_model_bl.load_weights('autoencoder_bl_weights.weights.h5',
                             skip_mismatch = True)
encoder_model_bl.compile(
    loss = hp_loss_fn,
```

```

        optimizer = hp_optimizer,
        metrics = hp_metrics,
    )

autoencoder_model_sm.save_weights('autoencoder_sm_weights.weights.h5',
                                  overwrite = True)

encoder_model_sm = keras.Model(inputs = input_data, outputs = Encoder_sm(input_data))
encoder_model_sm.load_weights('autoencoder_sm_weights.weights.h5',
                             skip_mismatch = True)
encoder_model_sm.compile(
    loss = hp_loss_fn,
    optimizer = hp_optimizer,
    metrics = hp_metrics,
)

autoencoder_model_bg.save_weights('autoencoder_bg_weights.weights.h5',
                                  overwrite = True)

encoder_model_bg = keras.Model(inputs = input_data, outputs = Encoder_bg(input_data))
encoder_model_bg.load_weights('autoencoder_bg_weights.weights.h5',
                             skip_mismatch = True)
encoder_model_bg.compile(
    loss = hp_loss_fn,
    optimizer = hp_optimizer,
    metrics = hp_metrics,
)

```

Let us embed our genotyping data using the encoder we have just constructed. We can also visualise the embedding.

```
In [11]: embedded_points_bl = encoder_model_sm.predict(geno_data.to_numpy())
print(embedded_points_sm)

x = embedded_points_sm[:,0]
y = embedded_points_sm[:,1]
pop = pheno2['population']
data = {'x':x, 'y':y, 'pop':pop}
plt.figure(figsize = (10,10))
sns.scatterplot(x='x', y='y', data=data, hue='pop', style='pop', s=100)
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0, markerscale=2)
plt.show()

embedded_points_sm = encoder_model_sm.predict(geno_data.to_numpy())
print(embedded_points_sm)

x = embedded_points_sm[:,0]
y = embedded_points_sm[:,1]
pop = pheno2['population']
data = {'x':x, 'y':y, 'pop':pop}
plt.figure(figsize = (10,10))
sns.scatterplot(x='x', y='y', data=data, hue='pop', style='pop', s=100)
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0, markerscale=2)
plt.show()

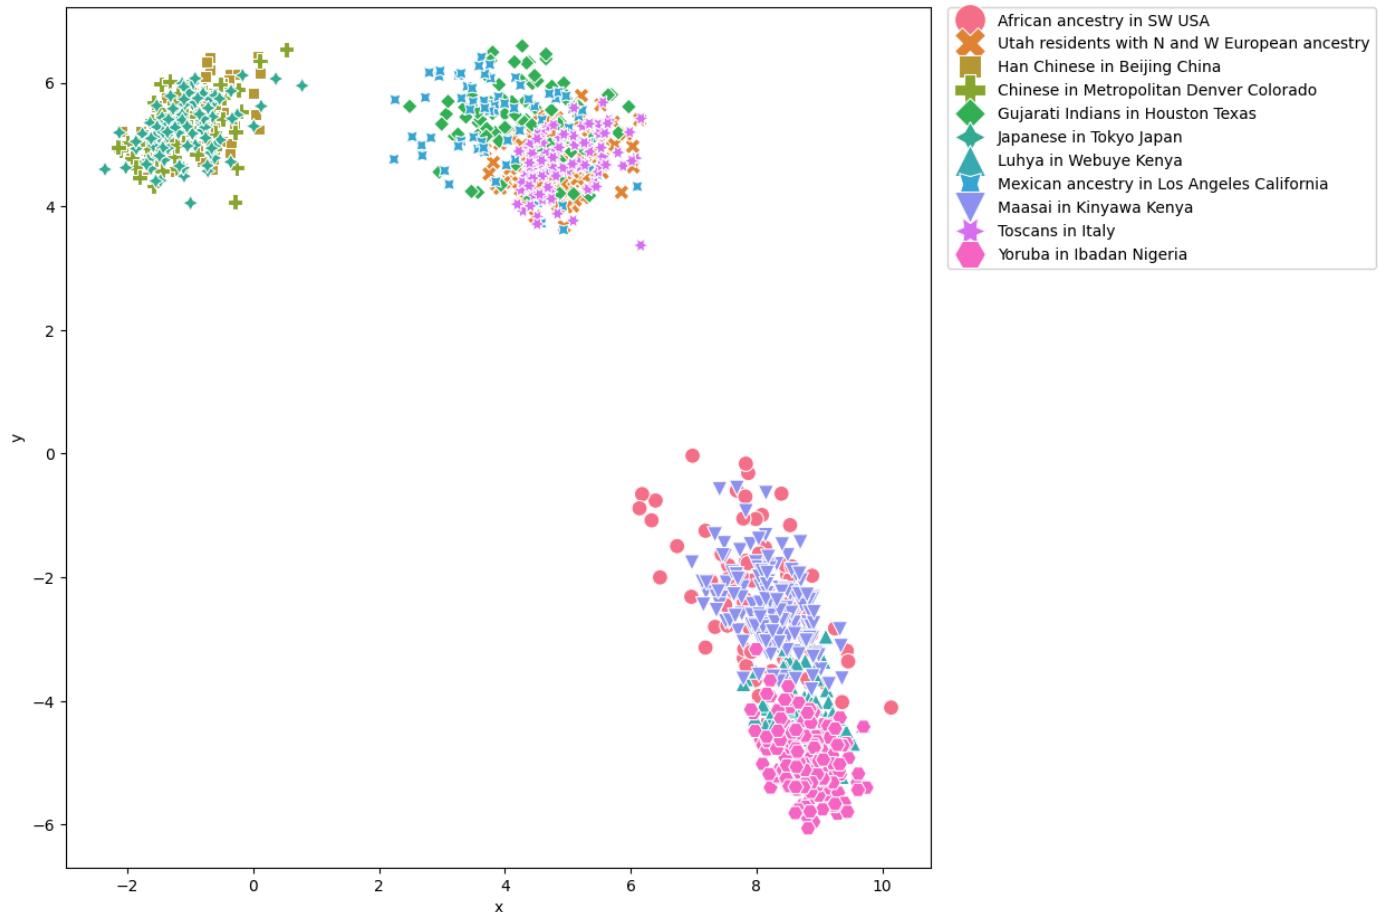
embedded_points_bg = encoder_model_bg.predict(geno_data.to_numpy())
print(embedded_points_bg)

x = embedded_points_bg[:,0]
y = embedded_points_bg[:,1]
pop = pheno2['population']
data = {'x':x, 'y':y, 'pop':pop}
plt.figure(figsize = (10,10))
sns.scatterplot(x='x', y='y', data=data, hue='pop', style='pop', s=100)
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0, markerscale=2)
plt.show()
```

2025-11-27 11:26:41.409475: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_4 0', 8 bytes spill stores, 8 bytes spill loads

37/37 ━━━━━━ 2s 4ms/step

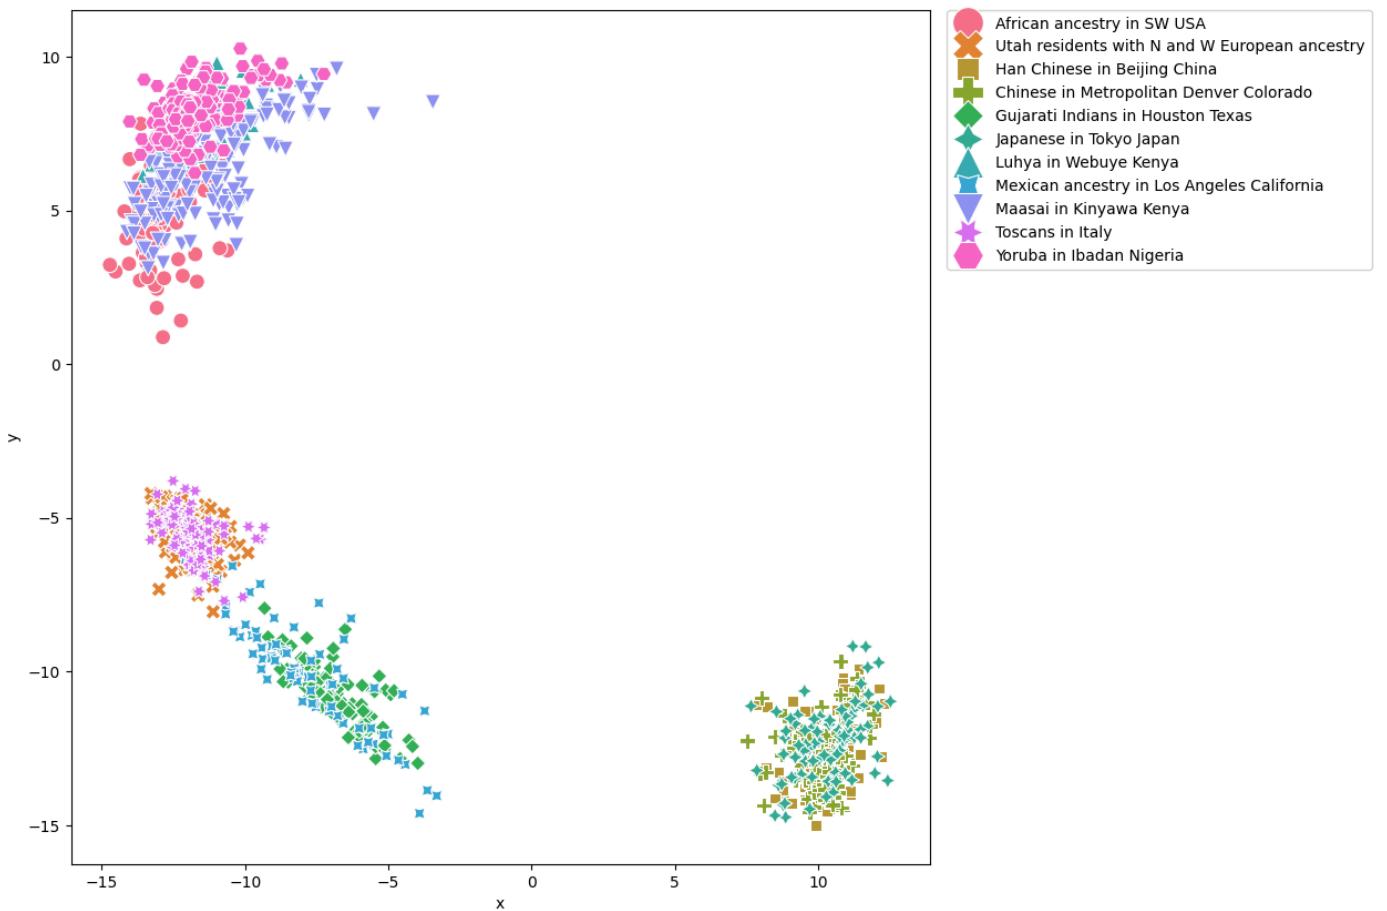
```
[[ 8.099391 -2.7050934]
 [ 8.437094 -3.3191807]
 [ 8.485687 -1.9565052]
 ...
 [ 8.338465 -4.4818306]
 [ 8.988413 -4.9713206]
 [ 9.060453 -4.947927 ]]
```



2025-11-27 11:26:43.566575: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_3 8', 4 bytes spill stores, 4 bytes spill loads

37/37 ━━━━━━ 1s 4ms/step

```
[ [-12.144369  7.4149075]
 [ -11.999359  6.5924506]
 [ -13.135323  5.2313647]
 ...
 [ -12.727489  7.248802 ]
 [ -11.991169  7.9087834]
 [ -10.5780325 8.285355 ]]
```



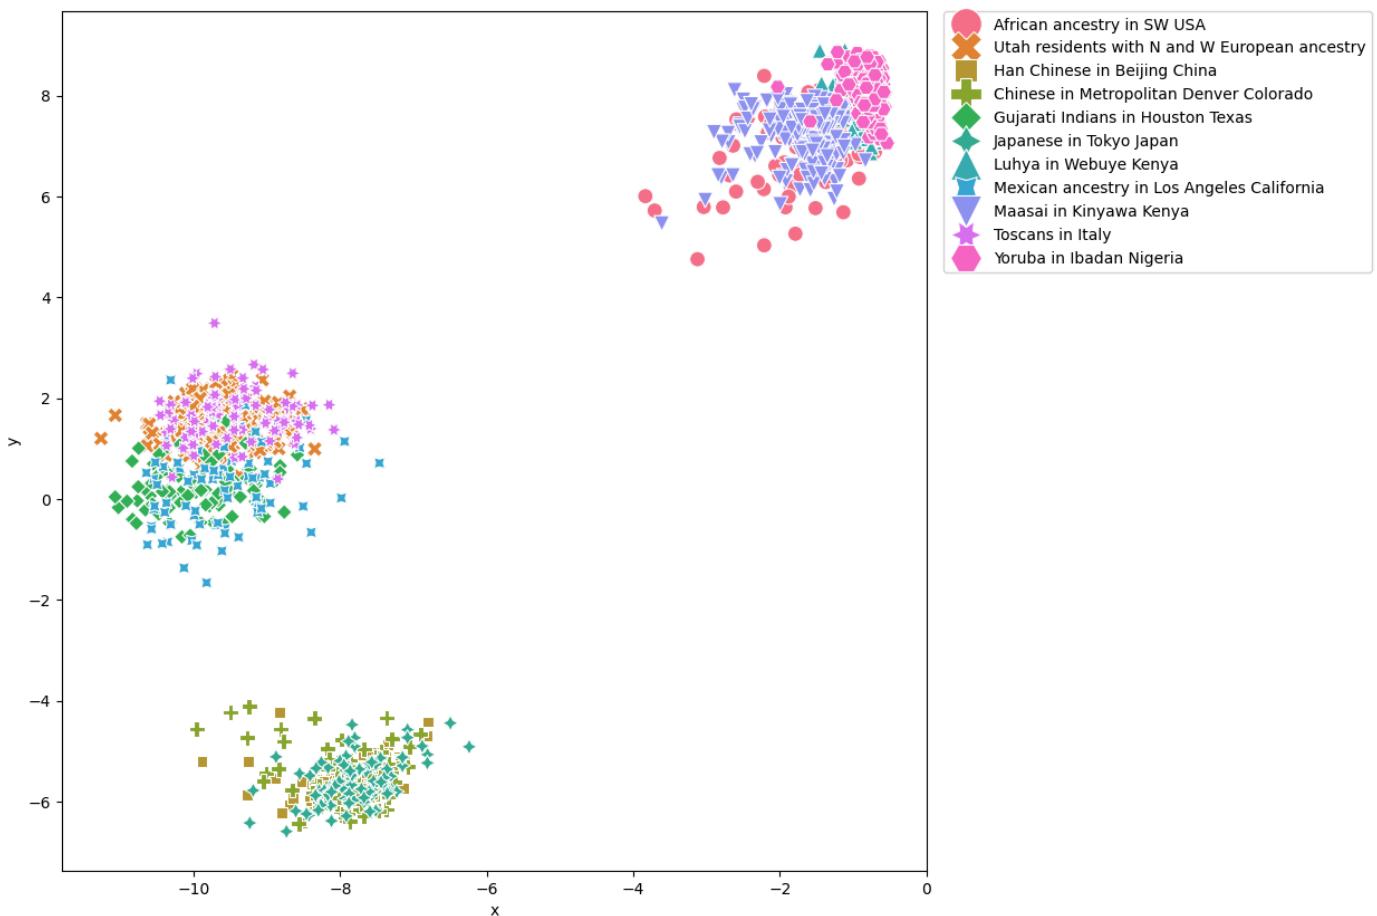
```
2025-11-27 11:26:46.553873: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_40', 160 bytes spill stores, 160 bytes spill loads
```

```
2025-11-27 11:26:46.700165: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_40', 356 bytes spill stores, 356 bytes spill loads
```

```
2025-11-27 11:26:47.275663: I external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:346] ptxas warning : Registers are spilled to local memory in function 'gemm_fusion_dot_40', 700 bytes spill stores, 700 bytes spill loads
```

37/37 ━━━━━━━━ 3s 4ms/step

```
[ [-1.7340165  7.460173  ]
 [ -0.8109024  7.1927357 ]
 [ -1.9854507  7.8065553 ]
 ...
 [ -1.0569018  8.219454  ]
 [ -0.84613174 7.8077846 ]
 [ -0.58026487 8.071715  ]]
```



Now, we will compare the result with:

- MDS on the kinship matrix
- PCA performed directly on raw genotypes

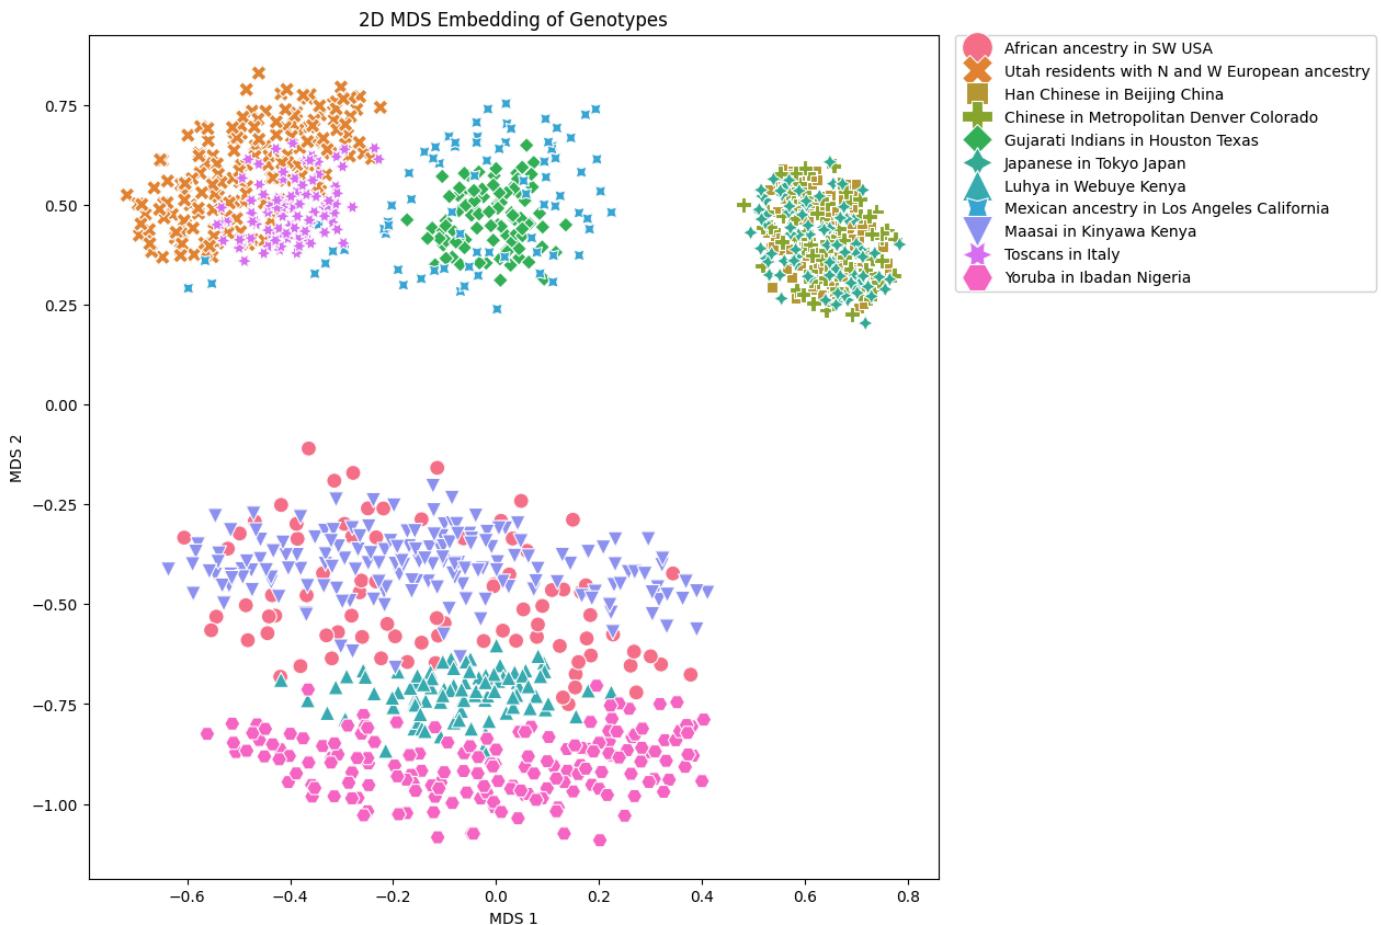
```
In [12]: kinship = pd.read_csv("autosomal_5k_kinship.csv", header=0, index_col=0)
```

```
In [ ]: from sklearn.manifold import MDS
embedding = MDS(n_components=2)
mds_embedding = embedding.fit_transform(kinship)
```

```
In [14]: # TASK Plot MDS embedding in a way similar to plotting autoencoder embeddings
x = mds_embedding[:,0]
y = mds_embedding[:,1]
data = {'x':x, 'y':y, 'pop':pop}

plot_df = pd.DataFrame({'x': x, 'y': y, 'pop': pop})

plt.figure(figsize = (10,10))
sns.scatterplot(x='x', y='y', data=plot_df, hue='pop', style='pop', s=100)
plt.xlabel('MDS 1')
plt.ylabel('MDS 2')
plt.title('2D MDS Embedding of Genotypes')
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0, markerscale=2)
plt.show()
```



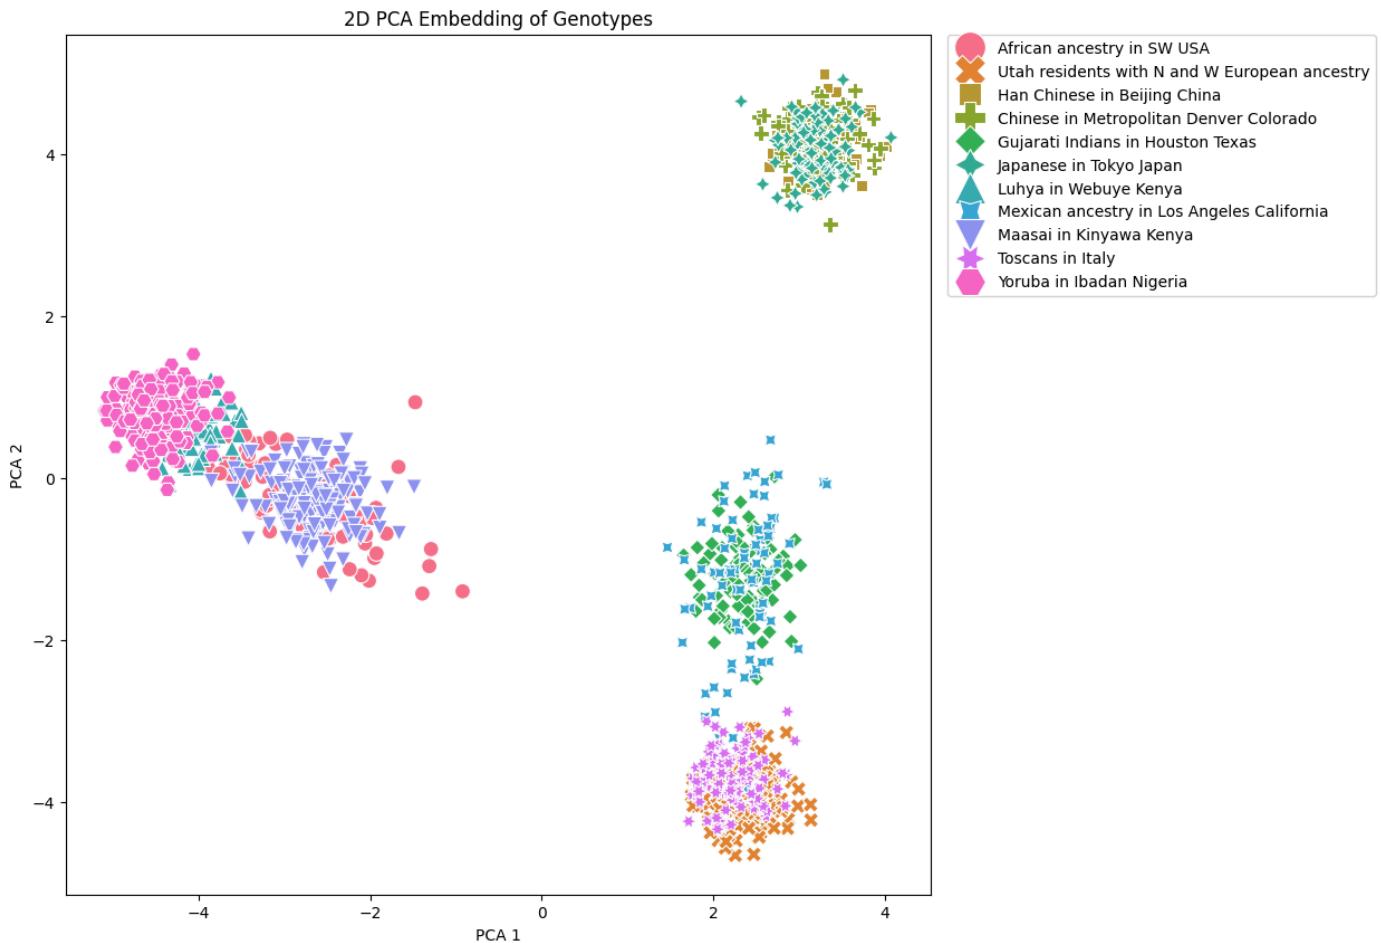
Finally, we will perform PCA on raw genotypes:

```
In [15]: # TASK perform PCA with 2 components on raw genotypes (use geno_data as input but remember it
# Visualise the result.
from sklearn.decomposition import PCA
embedding = PCA(n_components=2)
pca_embedding = embedding.fit_transform(geno_data.to_numpy())

x = pca_embedding[:,0]
y = pca_embedding[:,1]
pop = pheno2['population']

plot_df = pd.DataFrame({'x': x, 'y': y, 'pop': pop})

plt.figure(figsize=(10,10))
sns.scatterplot(x='x', y='y', data=plot_df, hue='pop', style='pop', s=100)
plt.xlabel('PCA 1')
plt.ylabel('PCA 2')
plt.title('2D PCA Embedding of Genotypes')
plt.legend(bbox_to_anchor=(1.02, 1), loc='upper left', borderaxespad=0, markerscale=2)
plt.show()
```



Final comments

- Task 1: Each locus in the dataset is a SNP (one column per locus in the genotype matrix). The first loci are: rs416967, rs17013842, rs13052452, rs11049986, rs10994341, rs1504289, rs882529, rs3885937, rs537330, rs9372090. These rsIDs are unique identifiers but do not show chromosome information. So, we cannot tell the chromosome of a SNP just from its name.
- Task 2: Code implementation.
- Task 3: The autoencoder gives the clearest separation of populations because it can capture non-linear relationships between markers. MDS works fairly well but has more overlap between genetically close groups. PCA, as a linear method, is less able to separate populations with high correlation or gene flow. In general, genetically similar populations overlap more, while the autoencoder is less sensitive to outliers and forms more compact clusters.
- Task 4: All three autoencoder architectures (Baseline, Smaller, Bigger) capture the relative structure of the populations. Clusters are compact and mostly well-separated, with Bigger having the tightest clusters. The absolute positions in 2D differ between models because of random initialization and arbitrary rotations or translations, which is normal. The main difference between architectures is the compactness of clusters and how well closely related populations are resolved, not the overall layout.